




Kubernetes: The Final Frontier

Amanda Kay Moran



#ossummit

@AmandaK_Data



The Final Frontier....

This is an introduction to Kubernetes tutorial.

The continuing mission: to explore the strange new worlds of microservices, containerization, and their management. To seek out new skills and new adventures. To boldly go where no one has gone before!



Agenda

- What is Kubernetes?
- Why is it so Popular?
- Why Should I learn about Kubernetes?
- Kubernetes Architecture
- What is a Pod? What is a Deployment?
- Install MiniKube
- Deploy Simple Application
- Deploy Web Application
- How to have High Availability and Scalability
- Where to Learn more!



Who is Amanda?

- Bay Area based Software Engineer/Solutions Architect
- Machine Learning, Analytics, Distributed Systems
- Variety of Companies, big and small!
- Apache Committer and PMC Member Apache Trafodion
- What Do I Love:
 - Dogs
 - Disneyland
 - Veggies
 - Teaching, training, helping others
 - Running and Exercise





Hands On Lab: Prework

- Install a Hypervisor
- Personally, I use [Virtualbox](#)
- 2 minutes to give folks the ability to kick off their download
- We will also be using [Minikube](#) --if you already have virtualbox installed



What is Kubernetes?

- Kubernetes (K8s) is an open-source system for automating deployment, scaling, and management of containerized applications.
- Donated to the CNCF foundation by Google and has been developed and used at Google for over 15 years
- First release was in June 2014 and is now 6 years old





What is Kubernetes?

- **Benefits of Kubernetes**

- Run Applications Anywhere
- Easy cluster management
- Service Discovery and load balancing
- Storage management
- Automated rollouts and rollbacks
- Automatic bin packing -- placing containers by resources
- Self Healing
- Horizontal Scaling



Why is it so Popular?

- Companies are moving away from monolithic applications

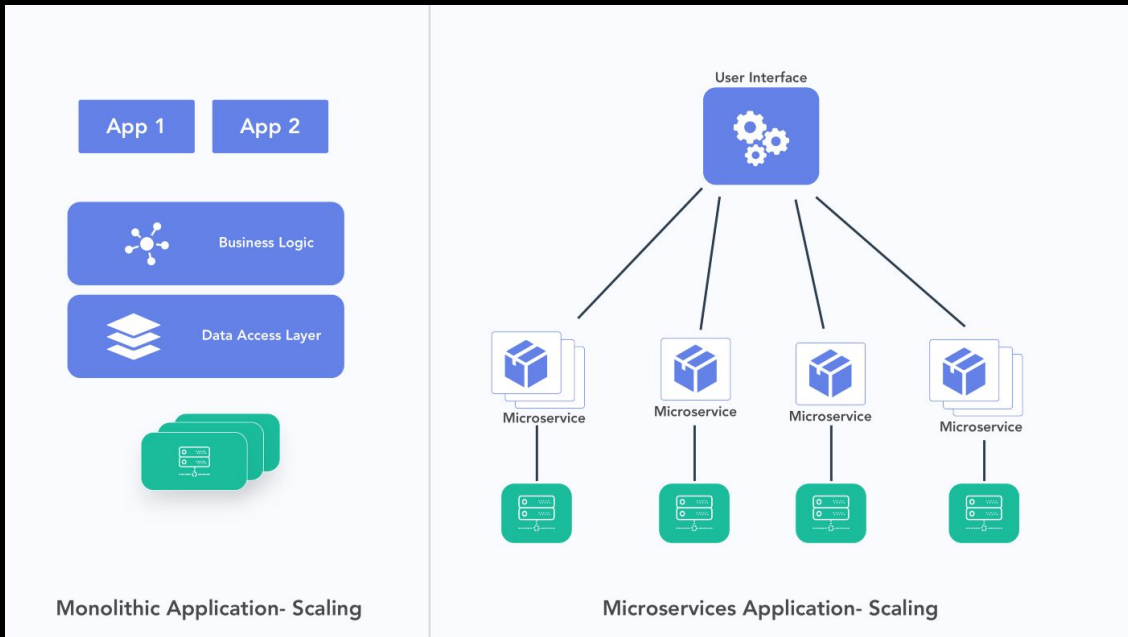



Photo Credit: [Microservices Orchestration with Kubernetes](#)



Why is it so Popular?


- Why use Microservices?
 - A microservice is a service that just does a single task and that is all
 - Rapid development
 - Ability to swap out components of an architecture with ease
 - Easy to automate for CI/CD
 - Flexibility → easier to change course
- Microservices are easy to containerize
 - Container Services: Docker
- Kubernetes is the best place to manage containers





Why is it so Popular?

- Many ways to install and run
- Manually
- Cloud Offering (Platform as a Service)
 - EKS
 - GKE
 - AKS
- Within minutes you can have a managed Kubernetes cluster
- 3 month release cycle





Why is it so Popular?


- Kubernetes has a very active and supportive community
- Many different company contribute to the project
- Special Interest Groups -- that meet regularly
- Meetups and Conferences

The screenshot shows the GitHub repository for Kubernetes. At the top, the repository name 'kubernetes / kubernetes' is displayed. To the right, there are buttons for 'Watch' (3.2k), 'Star' (66.9k), and 'Fork' (24k). Below these, a navigation bar includes links for 'Code', 'Issues' (2,032), 'Pull requests' (971), 'Actions', 'Projects' (8), 'Security' (0), and 'Insights'. The main heading reads 'Production-Grade Container Scheduling and Management' with a link to 'https://kubernetes.io'. Below this, there are tags for 'kubernetes', 'go', 'cncf', and 'containers'. At the bottom, a statistics bar shows: 91,631 commits, 41 branches, 0 packages, 654 releases, 2,558 contributors, and Apache-2.0 license. A progress bar is visible at the very bottom of the repository page.



Why Should I learn about Kubernetes?

- If you work in DevOps or Infrastructure
 - No Brainer!
- Very popular new technology
- But what about Developers and Data Scientist
 - The infrastructure will affect how build/use applications
 - Will affect how you build models
 - Storage
 - Hardcoding
 - Bringing in new package → not persistent
- An Introduction!





Kubernetes Architecture

- Terminology
- Control Plane
 - kube-api-server
 - kube-controller-manager
 - kube-scheduler
 - cloud-control-manager
 - Etcd (Key Value Database)
- Node
 - Kubelet
 - kube-proxy
 - Container Runtime



Kubernetes Architecture

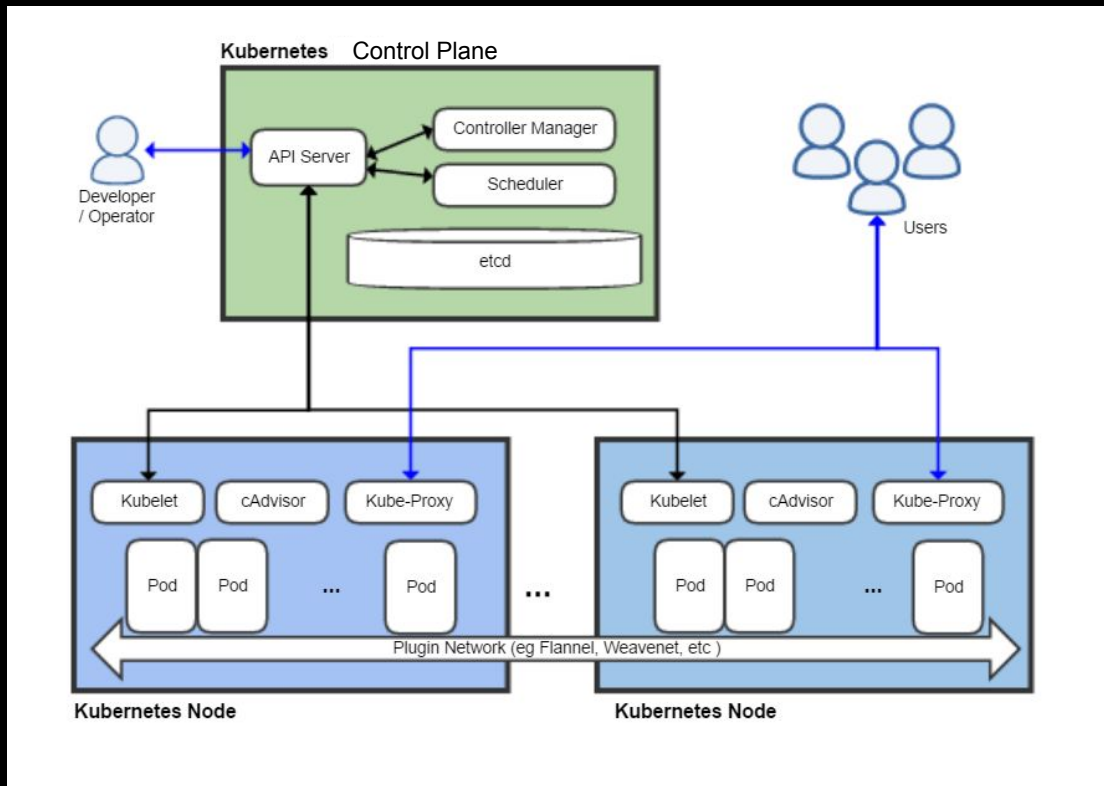


Photo Credit: [Wikipedia: Kubernetes](#)



Kubernetes Architecture

- Control Plane

- kube-api-server
- kube-controller-manager
- kube-scheduler
- cloud-control-manager
- Etcd (Key Value Database)

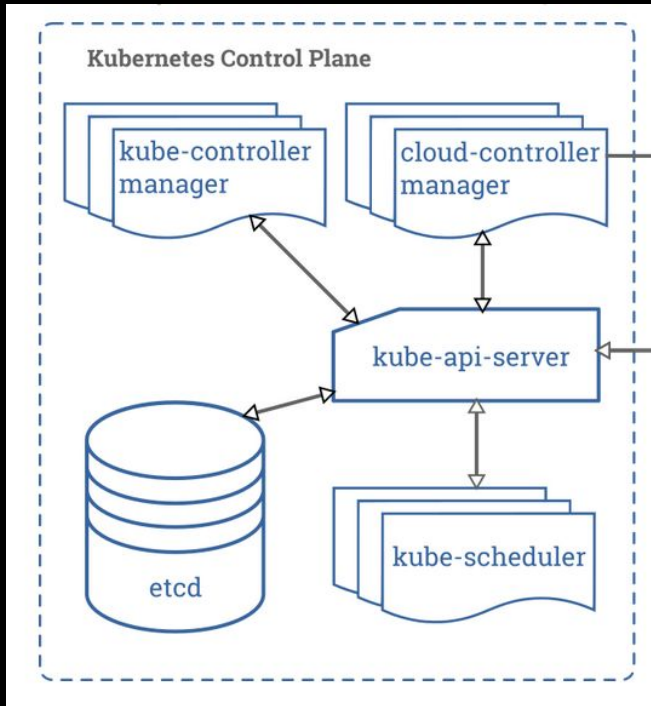
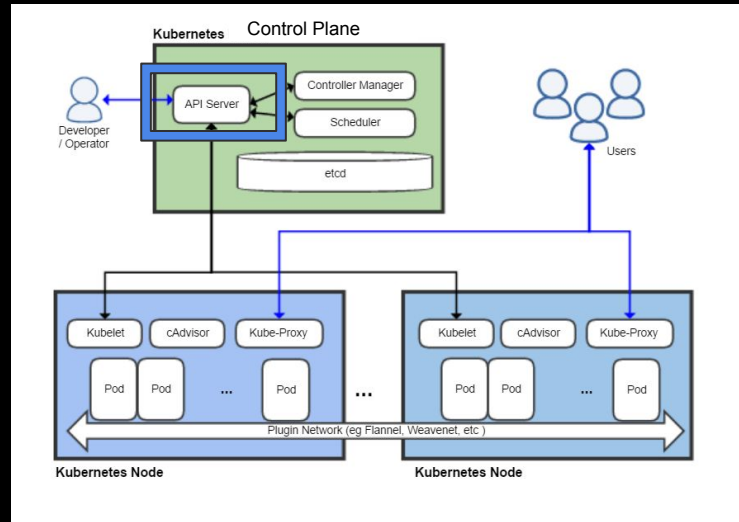


Photo Credit: [Kubernetes Documentation](#)



Kubernetes Architecture

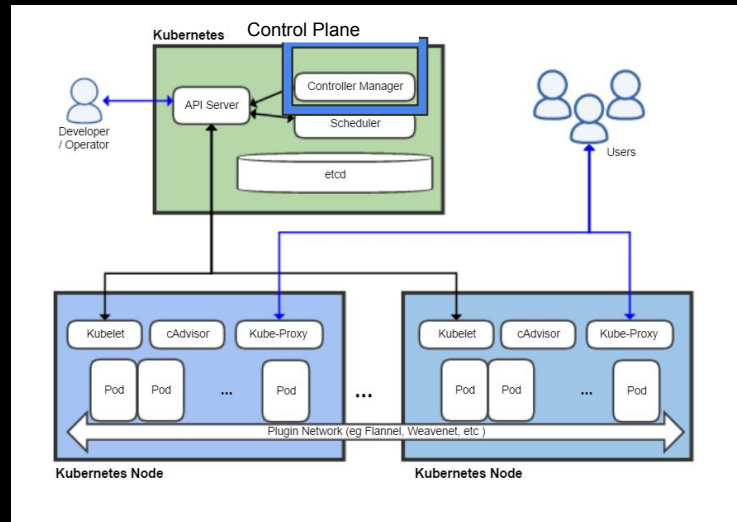
- kube-api-server
 - Responsible for the Kubernetes API
 - How you will interact with the k8s cluster
 - Will use a command line tool kubectl to interact
 - Interacts with etcd, scheduler, controllers



Kubernetes Architecture



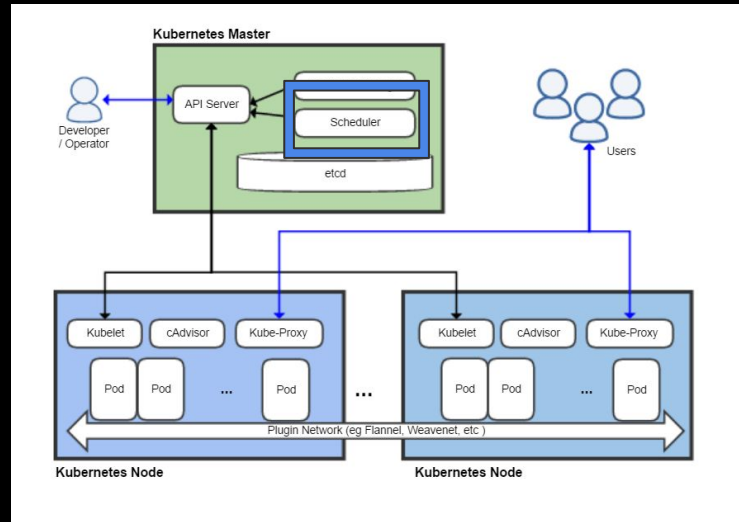
- kube-controller-manager
 - Runs multiple controllers
 - Node controller
 - Replication Controller
 - Service Account and Token Controller





Kubernetes Architecture

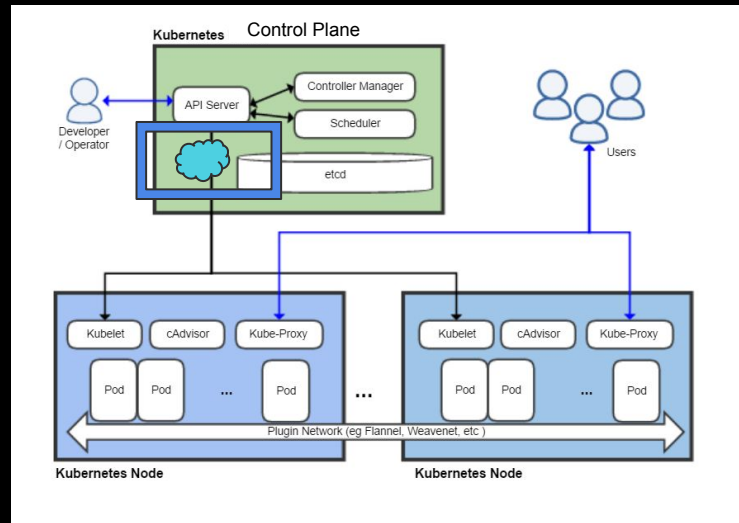
- kube-scheduler
 - Scheduler for the pods on the cluster
 - Resources needed
 - Affinity
 - Data locality
 - Uses different algorithms (configurable) to place
 - Supply and demand



Kubernetes Architecture



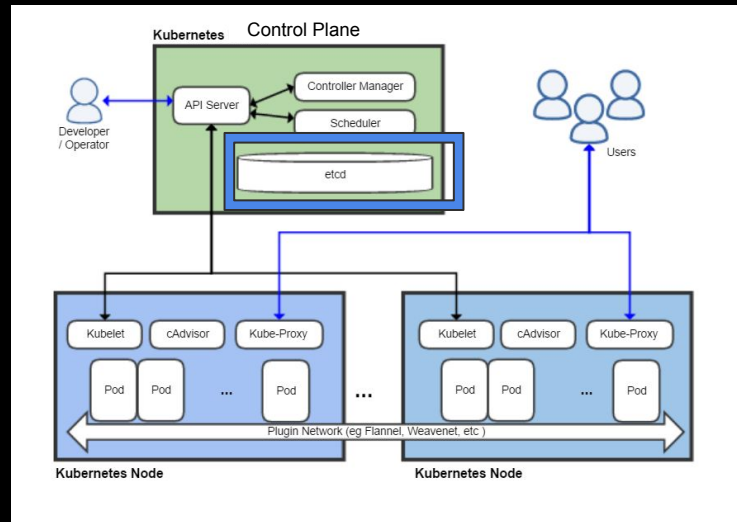
- cloud-control-manager
 - Ability to use your Cloud providers API
 - Only have in the cloud (not on perm or Minikube)



Kubernetes Architecture



- etcd
 - Key value consistent database
 - Consistent over Available → CAP Theorem
 - Stores all activity on the cluster
 - Combined with API Server to perform actions
 - The API Server used watch API on etcd to monitor





Kubernetes Architecture

- Node
 - Kubelet
 - kube-proxy
 - Container Runtime

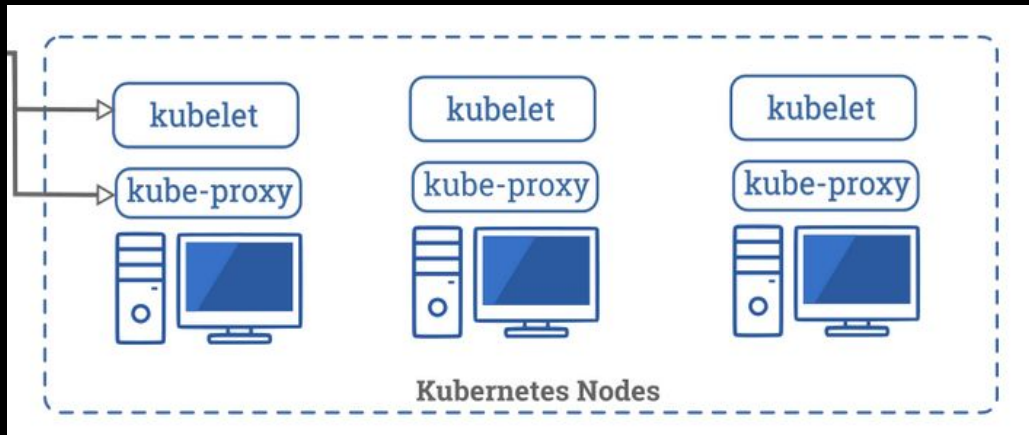


Photo Credit: [Kubernetes Documentation](#)



Kubernetes Architecture

- Kubelet
 - A process (agent) that runs on each kubernetes node
 - Uses Pod Specifications to understand which pods and containers should run
 - Monitors node it is responsible for
 - Control plane and kubelet work closely together

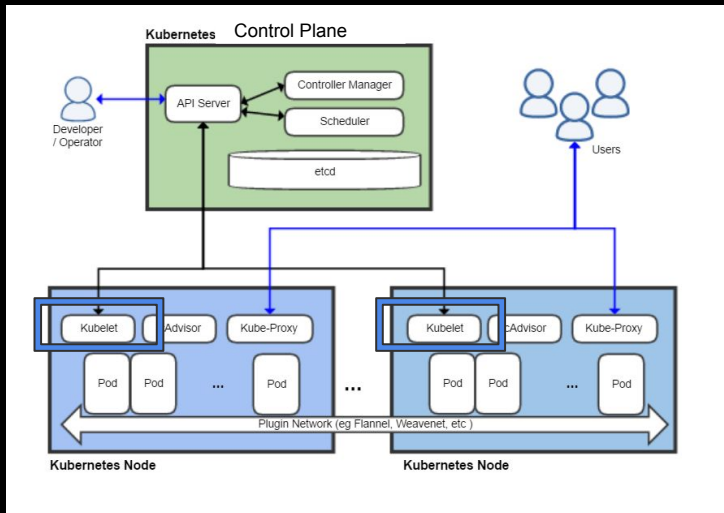


Photo Credit: [Kubernetes Documentation](#)



Kubernetes Architecture

- kube-proxy
 - Network proxy that runs on each node
 - Uses network rules to allow for pods to talk to each other inside and outside the cluster

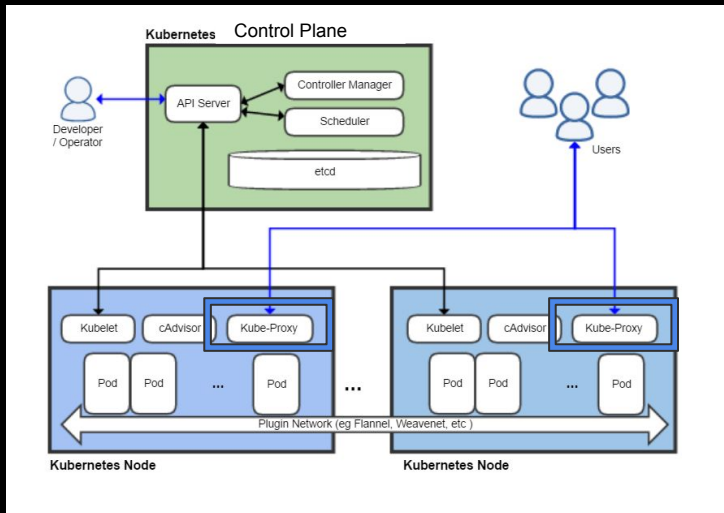


Photo Credit: [Kubernetes Documentation](#)



Kubernetes Architecture

- Container Runtime
 - Containers live within pods and package our application
 - Must be installed on each node
 - Doesn't have to just be Docker
 - Containerd
 - Crio

Workloads: What is a Pod?



Photo Credit: [Science News](#)




Photo Credit: [MatthewPalmer.net](#)



Workloads: What is a Pod?

- Smallest unit within managed Kubernetes
 - Containers run outside Kubernetes not managed
- Pods configured by a yaml file
- Will pull a container image from a source
- Containers in a pod are always colocated
- Unique internal ip address per pod
- Containers talk to each other via localhost
- All containers use common volume storage
- Managed by the kube api or controller

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   creationTimestamp: null
5   labels:
6     run: pod
7   name: pod
8 spec:
9   containers:
10    - image: amoran06/simple
11      name: pod
12      resources: {}
13    dnsPolicy: ClusterFirst
14    restartPolicy: Always
15 status: {}
```



Controllers: What is a Controller?

- A controller is a control loop that manages the environment
 - Never terminates and continues to monitor the situation
- kube-control-manager helps with managing each controller type
- It works on the current state and gets you to the desired state
 - I have two pods and I want 3
- Think of an Oven sensor





Controllers: What is a Deployment?

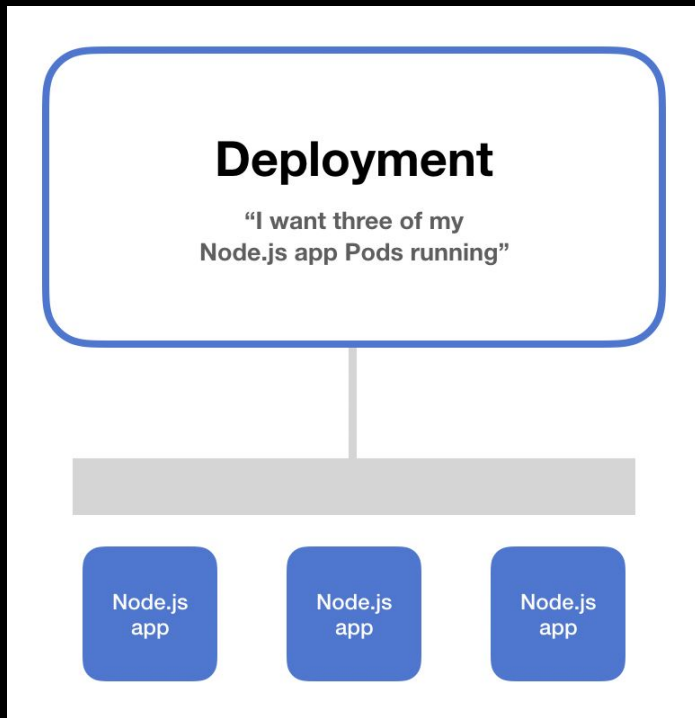


Photo Credit:
[MatthewPalmer.net](https://matthewpalmer.net)



Controllers: What is a Deployment?

- A deployment is setting a desired state for your pods
- Delete and Add Pods
- Add a ReplicaSet → How many I want
- Do I want a restart?
 - Batch job that just runs once
- Ability to easily update and upgrade

Lab Time



Go to: <https://github.com/amandamoran/opensourceSummit.git>


#ossummit

@AmandaK_Data



Continue Learning...

- Reading Kubernetes [Docs](#)
- Udacity course: [Introducing Scalable Microservices Kubernetes](#)
- [Kubernetes in Action](#) by Marko Luksa
- Linux Foundation Training: [Introduction to Kubernetes](#)





Become Apart of the Community!

- Join a Meetup
- Make a Pull Request
- Keep attending Conferences
- Answer Questions
- Write docs
- Everyone is welcome!



References

- Kubernetes [Docs](#)
- Kubernetes [Wikipedia](#)

Thank you!





#ossSummit

@AmandaK_Data